Discussion 11

Caches

Aditya Balasubramanian

`aditbala [at] berkeley [dot] edu`

Announcements

Agenda

- Cache Basics
- TIO Breakdown
- Cache Workflow
- 3 C's of Cache Misses
- Multilevel Caches + AMAT

Caches

4

Caches

- The closer we get to the CPU, the faster we can load data
- Registers
  - On the CPU, blazingly fast, super tiny
- Memory
  - On the motherboard, still pretty fast, pretty big
- Between them, we have the Cache
  - On the CPU, faster than memory, small-ish
- We turn memory addresses into TIO bits to place it into the cache

Data that's close

- Spatial Locality
  - Close with respect to location in memory
  - Data that "lives" +/- a certain amount from the accessed address might be accessed soon
- Temporal Locality
  - Close with respect to time
  - Data that's accessed recently might be accessed soon again

Cache Terminology

- Cache line
  - a chunk of data that is loaded into the cache from main memory upon an access (also known as block)
- Set
  - groups of cache lines to form a larger group indexed by the index
- Associativity
  - number of cache lines needed to form a set (4-way associative cache means each set contains 4 blocks)

7

Types of Caches

- Direct Mapped
  - Each address can only go into one fixed cache line
  - Easiest to implement on hardware
- N-way Set Associative
  - Each address can go into one fixed set (N cache lines)
- Fully Associative
  - Each address can go into any cache line

TIO Breakdown

tag	index	offset

- Tag
  - $\log_2$ (Memory size in bytes) - I - O
- Index
  - $\log_2$ (Number of blocks / associativity) = log2(Number of indices)
- Offset
  - $\log_2$ (Size of block in bytes)

Cache Workflow

10

Workflow

- **We know we're accessing address, say, 0xDEADBEEF, now what?**

1. **Breakdown address** based on TIO

2. **Use index** to figure out which set our cache line goes into

3. Within the set, **check the tag** of every existing cache line to see if the given tag matches.
   a. If any tag matches, go to step 4, otherwise go to step 5

4. Tag match! **Check the valid bit** to make sure we still have an updated copy of the cache line.
   a. If the valid bit is 1, then data is already in the cache, so we can get the data using the offset
   b. If valid bit is 0 (invalid), continue with step 5

5. **Fetch the entire cache line** (block aligned!) into the cache
   a. What if the set is full?
   i. We need to evict some previous line…
   ii. Which line to evict?

Eviction Policy

- Least Recently Used (LRU)
  - Probably most common
  - Least recently used block get evicted
- First In First Out (FIFO)
  - The oldest block in the set gets evicted
- Others
  - Most Recently Used (MRU)
  - First In Last Out (FILO)

3 C's of Cache Misses

Cache Misses

- Compulsory Miss
  - We haven't seen this block before
  - ALWAYS happens the first time we pull a block into the cache
  - Can be reduced by increasing block size
- Conflict Miss
  - We had to evict this block because some other block replaced it
  - Can be reduced by increasing associativity or having a better replacement policy
- Capacity Miss
  - We had to evict this block because the cache wasn't big enough
  - Can be reduced by increasing cache size

Multi-level Caches

- Usually has L1, L2 cache and occasionally L3
- Only cache misses from lower level will go to the next level!
  - If L1 hits, no need to look in L2
- Global Miss Rate
  - Miss rate of a cache level in the scope of the entire system
  - MR = # misses at this level / # total accesses of cache
- Local Miss Rate
  - Miss rate of a cache level in the scope of the level itself
  - MR = # misses at this level / # total access in this level

AMAT

`AMAT = Hit Time + Miss Rate x Miss Penalty`

- Hit Time
  - Time it takes to get the data if it's in the cache already
- Miss Rate: local miss rate
  - # misses / # accesses in a level
- Miss Penalty
  - Time it takes to get the data if it's not in the cache
  - Single-level cache: MP = time it takes to access main memory
  - Multi-level cache: MP = AMAT of the next level
- e.g. 2 level cache

`AMAT = L1 HT + L1 MR * (L2 HT + L2 MR * Main Memory Access)`

Thank you!